
Gaussian Evolution Strategies

Bram Pulles

Department of Artificial Intelligence
Radboud University
S1015194
bram.pulles@ru.nl

Abstract

The inverted pendulum problem is a widely known simple environment in which different machine learning approaches are tested. We compare the performance of the reinforce algorithm and a simple Gaussian evolutionary strategy (ES) for solving the inverted pendulum problem. We rely heavily on hardware acceleration using JAX and show that the ES implementation performs a magnitude faster than the reinforce implementation. We conclude that the ES is a promising approach with good convergence properties for this particular problem.

1 Introduction

Reinforcement learning (RL) constitutes a broad field of machine learning methods in which agents are trained by rewarding desirable behaviours and punishing undesirable behaviours. We aim to teach an agent how to solve the inverted pendulum problem. In this problem the agent guides the bob (the object on the pendulum) to the unstable fixed point at the top of the pendulum, see figure 1. Using RL methods the agent can be trained to show this desired behaviour.



Figure 1: Pendulum with the bob at the unstable fixed point.

We have previously shown how the reinforce algorithm, also known as the Monte Carlo policy gradient algorithm, can be utilized to teach the agent how to solve the inverted pendulum problem.¹ This method is based on the idea of using gradient descent to optimize a policy by directly optimizing the cumulative reward. Instead of using a gradient-based RL method, we contrast the approach with a gradient-free simple Gaussian evolution strategy.

Evolutionary strategies (ES) and gradient-based approaches are two different families of optimization methods each with their own strengths and weaknesses. Here are some reasons why ES might be preferred over gradient-based approaches:

¹This was done in assignment 8. We use the provided solution as reference implementation.

- ES can perform well in stochastic environments. Since ES samples solutions from a distribution it is flexible and can adapt to uncertain environments better than gradient-based methods which might get stuck in a local solution space due to the noise. In general, ES tend to explore the search space more globally than gradient-based methods due to this property.
- ES can be applied to problems where gradients may not be available or expensive to compute. This includes problems with discontinuous and non-differentiable objective functions. It would be challenging to use gradient-based methods for solving such problems.
- ES are easy to parallelize. ES can exploit distributed hardware or other parallel computing power. This way, the performance, or cumulative reward in our case, of various solutions can be computed simultaneously. This can lead to faster convergence times, especially when taken advantage of hardware capabilities with a library like JAX.

It is clear now why ES are interesting when contrasted to other gradient-based approaches. In section 2 we define a research question and outline the ES that we implement to solve the inverted pendulum problem. In section 3 we discuss the performance of the ES and compare this to the previously implemented reinforce algorithm. We end with a discussion of the results in section 4.

2 Methods

2.1 Research question

In line with the introductory goals, we aim to train an agent for solving the inverted pendulum problem. To this end we use the pendulum implementation from Gym, a toolkit for developing and comparing reinforcement learning algorithms made by OpenAI. We implement a straightforward Gaussian evolutionary strategy. For both the ES and the reinforce algorithm we aim to answer the following question:

How does the simple Gaussian ES perform compared to the reinforce algorithm in terms of convergence, measured in time and in number of iterations needed?

By considering the runtime, we benchmark the training duration for each algorithm. Additionally, the number of iterations needed by either algorithm to converge provides another metric giving an impression of the performance. We expect that the ES will converge faster compared to the reinforce algorithm, due to the high cost of calculating the cumulative reward for the objective function and the parallelisability of the ES. For the implementation we use JAX which can make good use of the available concurrency capable hardware.

2.2 Evolutionary strategy

The agents that we use for both the reinforce algorithm and the ES consist of a simple neural network with one hidden layer. The network is supposed to apply a force to the bob such that the bob is balance in the unstable fixed point. To this end it gets information at every time step. As input it gets the (x, y) position and the angular velocity of the bob. As output it gives a probability distribution over three possible actions: -1, 0 or 1 torque applied to the free end of the pendulum. The agents performance is measured by computing the mean cumulative reward across multiple rollouts.

The neural network completely defines the agent. Let all the parameters, i.e. the weights and biases for each layer, be defined in one long vector θ . Our goal is to find an optimal θ such that the inverted pendulum problem is solved by the neural network defined as $\bar{\theta}$.

In order to do this using the evolutionary strategy we start with creating a n -dimensional Gaussian distribution $p_{\theta} \sim \mathcal{N}(\bar{\mu}, \bar{\sigma}^2 I)$ with $\bar{\mu}$ as the best estimate for each parameter of $\bar{\theta}$ and $\bar{\sigma}$ the uncertainty for each parameter. The function p_{θ} can be initialized with some vector $\bar{\mu}$ and $\bar{\sigma}$.

Next, we generate a population of N sample neural networks D from p_{θ} . This can be seen as the offspring from the current best estimate of $\theta = \bar{\mu}$ modelled in p_{θ} . For all the networks in D we compute the mean cumulative reward across multiple rollouts, giving us a performance score for each network in the population. We select the top $k < N$ performing sample networks and call this the elite set $D_{\text{elite}} \subset D$. Now we estimate the new variance as the squared difference between

all samples in the elite set and the previous mean $\bar{\sigma} = \frac{1}{N} \sum_{i=1}^k (\bar{\theta}_i - \bar{\mu})^2$. The new mean $\bar{\mu}$ is computed by taking the mean over all $\theta_i \in D_{\text{elite}}$. This leaves us with a new and hopefully better p_{θ} . We now throw away the old D and repeat these steps until $\theta = \bar{\mu}$ is satisfactory.

3 Results

The results for the training session of the reinforce algorithm are shown in figure 2. We use the provided hyperparameter settings: 64 hidden nodes, a learning rate of 0.005, gamma (the reward discount factor) is 0.99, the number of batches is 50 and the number of steps per episode is 200. The first time running this takes 82.6 seconds. Running it for the second time, which is faster because JAX does not need to compile again, takes 78.8 seconds.

The results for the training session of the Gaussian ES are shown in figure 3. The hyperparameter settings used are as follows: 10 hidden nodes, gamma is 0.99, the number of steps per episode is 200, we use 50 rollouts per sample to measure the performance, the population size N is 100 and we select the top $k = 10$ best performing networks, we start with $\bar{\mu} = \bar{0}$ (all zeros) and $\bar{\sigma} = 100I$ (a matrix with 100's on the diagonal). The first time running takes a mere 4.6 seconds and a second time only 3.1 seconds.

Given these results we can already conclude that the ES implementation is time wise a whole magnitude faster than the reinforce implementation. We also see that it takes only about 50 iterations to converge, compared to about 2000 iterations needed for the reinforce algorithm. Interestingly, when the ES implementation is run without being compiled through JAX, it takes over an hour to run, so JAX greatly decreases the runtime. We also found that for the Gaussian ES the choice of the initial $\bar{\sigma}$ matters a lot, if this is chosen to small the algorithm does not converge to a solution. The other hyperparameters have less influence on the training trajectory.




Figure 2: Reinforce algorithm training results.

4 Discussion

The results show a clear victory for the Gaussian ES compared to the reinforce algorithm, however there are a lot of side notes that should be placed here. First, we did not do any hyperparameter search. For the ES we have done this manually, while for the reinforce method we just took the given parameters. In order to do a fairer comparison a hyperparameter search is needed. Second, the reinforce implementation can likely be sped up more using JAX. For example, a forward pass through the network is not compiled now, while it is for ES method. Working more on the implementation details particularly with speed in mind can therefore improve the performance of both algorithms,




Figure 3: Gaussian evolutionary strategy training results.

but especially for the reinforce version. Third, the ES implementation can make good use of parallel hardware available in the computer used to run the training. If this would not be available the time performance of the ES would probably be a lot worse, since it heavily relies on parallelisation for computing a lot of forward passes through all the networks simultaneously.

While benchmarking the Gaussian ES there were quite some difficulties with finding the right hyperparameter settings. In particular because the importance of $\bar{\sigma}$ was overseen. There are two ways to tackle this problem. One is the already mentioned parameter search, this would likely find out that the value for $\bar{\sigma}$ is very important. Another solution would be to alter the ES currently used. As it is implemented now the variance depends heavily on the previous parameters of $\bar{\mu}$. An alternative method which aims to address to problems is the covariance matrix adaptation ES (CMA-ES). This algorithm could be implemented instead, but greatly increases the complexity.

Another important aspect of the analysis done is the problem that was chosen. In our case we tackled the inverted pendulum problem. For this problem the ES performed quite well. It would be interesting to expand the number of problems and compare against multiple problems to give a better comparison between the two methods.

All in all, the Gaussian evolutionary strategy gave us promising results. However, a more elaborate research is definitely required to provide us with a comparison that is trustworthy and representative of the performance differences between the Gaussian ES and the reinforce algorithm.

Acknowledgements

The simple Gaussian evolutionary strategy approach is taken from Lilian Weng’s blog.² The alternative method of CMA-ES was also found here. We also used the solution provided for assignment 8 and of course the syllabus written for the course. Lastly, we heavily relied on JAX in order to create a speedy implementation for both the algorithms.

²<https://lilianweng.github.io/posts/2019-09-05-evolution-strategies/>